Reproduction and Improvement on the BugsInPy Dataset

Faustino Aguilar
Dept. of Computer Engineering
University of Panama
Panama City, Panama
orcid.org/0009-0000-1375-1143

Samuel Grayson
Dept. of Computer Science
University of Illinois Urbana-Champaign
Urbana, IL
orcid.org/0000-0001-5411-356X

Darko Marinov
Dept. of Computer Science
University of Illinois Urbana-Champaign
Urbana, IL
orcid.org/0000-0001-5023-3492

Abstract—We assess the reproducibility of the BugsInPy dataset less than three years after its original publication. The bug dataset provides some information about the software environment in which the code should be run, but this information can be incomplete or can decay into something uninstallable over time. We rectify as many of these problems as we can and redesign the original dataset to be more easily reusable and reproducible by future research projects. Based on our experience, we offer suggestions to authors of Python artifacts to improve their reproducibility.

Index Terms—reproducibility, bug database, BugsInPy, Python, package managers, Pip, Conda, containers, Docker

I. INTRODUCTION

BugsInPy [1] is a curated dataset of real-world bugs in large Python projects, intended to be used by researchers to develop and evaluate software testing and debugging tools for Python on a diverse set of real-world bugs from multiple projects. This dataset can evaluate the efficacy of tools in bug detection, fixing, and software reliability. Many software engineering studies [2]–[6] use BugsInPy.

The BugsInPy dataset contains a variety of information about each bug, and these bugs are organized by the library they come from, including:

- A buggy commit
- A fixed commit
- Python version used
- Test cases that indicate the bug's presence

BugsInPy includes a database abstraction layer and a test execution framework. The database abstraction layer provides a way to access the dataset in a structured way. The test execution framework allows researchers to run tests relevant to a particular bug.

We sought to use the BugsInPy framework to verify that the test cases could be set up, that the buggy commit fails, and that the fixed commit passes. This work is a *reproduction* (using ACM's 2020 definition [7]) since we use the original work's scripts.

Our contributions are:

- Improvements to the test execution framework, which make it easier to run experiments *en masse*.
- Modifications to the BugsInPy test execution framework, which installs and uses the correct version of Python.
- The results of which bugs were reproducible.

We propose the following research questions:

RQ1. How many bugs in BugsInPy are reproducible with no "extra" work? For a bug to be reproducible, the software environment should install without failure, the buggy version should fail the identified test case, and the fixed version should pass.

RQ2. How many non-reproducible bugs can we rescue? We rescue a bug by modifying the scripts and data such that the bug is reproducible.

This article proceeds with the methodology section, which explains how we tried to reproduce BugsInPy and what rescue procedures we took when bugs were not reproducible. Then we summarize the results of our executions and analyze the failures. Finally, we engage in an open-ended discussion of our experiment with advice to authors of reproducible artifacts and those seeking to reproduce artifacts.

II. METHODOLOGY

As part of our rescue, we made the following changes:

1. We use a container build script to provide a consistent starting point in which our scripts will install the correct software environment. The container also sandboxes modifications the BugsInPy script wants to make (e.g., modifying `~/.bashrc`). While this image is available in this registry, we suggest that users seeking robust reproducibility build the image from our source rather than depending on an external registry.
2. For each bug, The BugsInPy script ignores the specified version of Python, deferring to the system default Python instead. Presumably, the BugsInPy authors manually changed their system's version of Python according to the specification of each bug, but this is not an automated process, making it difficult for future users. We modified this script to install the correct version of Python using

Conda. Conda is a cross-platform package manager. Packages installed by Conda neither use nor modify the system version of those packages, so Conda can support different environments, each with its own potentially conflicting requirements. Conda package repositories store packages containing pre-compiled binaries and metadata for each platform, so installing is much faster than compiling from source code.

3. The original BugsInPy scripts install all Python packages using Pip package manager. Pip can invoke the compiler to build dependencies from source code [8] or download prebuilt binary files. However, some packages may require additional system libraries or dependencies that cannot be installed solely through pip. For example, Matplotlib, a popular Python plotting library, has required system-level dependencies that pip cannot automatically handle, such as libpng, freetype, or Tk [9]. Consequently, if a bug in a project's environment depends on Matplotlib, attempting to install and run that project on a vanilla Ubuntu or Debian system without the necessary system libraries would result in installation failures. Presumably, the original authors manually modified their system to have these system libraries; in our case, we identify packages that Pip cannot install on vanilla Ubuntu or Debian and simply install those with Conda instead.

4. Building the environment from source code can be costly, so we reuse environments when the Python package requirements and Python versions are identical. This optimization helps reduce the time and resources required for environment setup, as it bypasses the costly process of building environments from source code. While it is tempting to use the same Conda environment for each project, there are multiple occasions where different bugs of the same project require different dependencies. For example, `ansible/bugs/{1,11,14}/requirements.txt` all vary subtly.

5. The BugsInPy framework correctly recognizes that installing the dependencies line by line `cat requirements.txt | filter | xargs -n 1 pip install`, rather than using `pip install -r requirements.txt`, bypasses certain restrictions imposed by pip. Specifically, when installing all dependencies at once, pip may ignore very old packages. However, sequentially installing the dependencies allows us to reproduce the bugs accurately. However, this results in failed installations for projects that include the `-e git+https://...` syntax in their `requirements.txt` file, since they would get passed along as `pip install -e` and `pip install git+https://...`. This revised code ensures that each line from the requirements.txt file is properly processed and passed as an argument to the `pip install` command. To address this issue, we have opened a pull request in the original repository [10]. This fix is crucial, as it impacts the reproducibility of

bugs in projects such as black, cookiecutter, keras, luigi, pandas, sanic, and thefuck.

III. RESULTS

TABLE I
REPRODUCTION OF BUGS IN BUGSINPY WITHOUT MODIFICATION

Project	Err	Bth-pass	Bth-fail	Exp	Total
PySnooper	2 (67%)	0 (0%)	0 (0%)	1 (33%)	3 (100%)
ansible	3 (17%)	0 (0%)	0 (0%)	15 (83%)	18 (100%)
black	1 (4%)	0 (0%)	0 (0%)	22 (96%)	23 (100%)
cookiecutter	2 (50%)	0 (0%)	0 (0%)	2 (50%)	4 (100%)
fastapi	0 (0%)	0 (0%)	0 (0%)	16 (100%)	16 (100%)
httpie	4 (80%)	0 (0%)	0 (0%)	1 (20%)	5 (100%)
keras	14 (31%)	0 (0%)	0 (0%)	31 (69%)	45 (100%)
luigi	33 (100%)	0 (0%)	0 (0%)	0 (0%)	33 (100%)
matplotlib	29 (97%)	0 (0%)	0 (0%)	1 (3%)	30 (100%)
pandas	47 (28%)	0 (0%)	0 (0%)	122 (72%)	169 (100%)
sanic	5 (100%)	0 (0%)	0 (0%)	0 (0%)	5 (100%)
scrapy	11 (28%)	0 (0%)	0 (0%)	29 (72%)	40 (100%)
spacy	2 (20%)	0 (0%)	0 (0%)	8 (80%)	10 (100%)
thefuck	8 (25%)	0 (0%)	0 (0%)	24 (75%)	32 (100%)
tornado	1 (6%)	0 (0%)	0 (0%)	15 (94%)	16 (100%)
tqdm	2 (22%)	0 (0%)	0 (0%)	7 (78%)	9 (100%)
youtube-dl	0 (0%)	0 (0%)	0 (0%)	43 (100%)	43 (100%)
Total	164 (33%)	0 (0%)	0 (0%)	337 (67%)	501 (100%)

> **RQ1.** We can reproduce 67% of the expected results in the unmodified BugsInPy dataset.

In Table I, the results we can get are:

- **Error**: Some part of the installation of the software environment needed to test the bug failed.
- **Both-pass**: Both versions pass; we would expect the buggy version to fail.
- **Both-fail**: Both versions fail; we would expect the fixed version to pass.
- **Expected**: The buggy version fails, and the fixed version passes. We consider *only* these bugs, "reproduced".

In that table, for each project, we show the raw count and percentage of outcomes for all bugs in that project. The last row shows the raw count and percentage of outcomes for all bugs in the dataset.

> **RQ2.** We were able to rescue 85% of the non-reproducible cases in the original BugsInPy, resulting in a total reproduction rate of 95%.

With over 95% of bugs being successfully reproduced and passing the tests, researchers have more bugs at their disposal for evaluating fuzzing, automatic program repair, and other research techniques.

Table III presents the running time required to reproduce bugs in each project within the BugsInPy dataset and run the respective containers. The provided running times are specific to the reproduction process on the given VM configuration, which had 8GB of RAM, 4 cores, and 100GB of free disk space. Reproduction times can vary depending on hardware resources, system configurations, and other environmental factors. The

Project	Err	Bth-pass	Bth-fail	Exp	Total
PySnooper	1 (33%)	0 (0%)	1 (33%)	1 (33%)	3 (100%)
ansible	0 (0%)	0 (0%)	0 (0%)	18 (100%)	18 (100%)
black	0 (0%)	0 (0%)	1 (4%)	22 (96%)	23 (100%)
cookiecutter	0 (0%)	0 (0%)	0 (0%)	4 (100%)	4 (100%)
fastapi	0 (0%)	0 (0%)	0 (0%)	16 (100%)	16 (100%)
httpie	0 (0%)	0 (0%)	0 (0%)	5 (100%)	5 (100%)
keras	3 (7%)	0 (0%)	1 (2%)	41 (91%)	45 (100%)
luigi	0 (0%)	6 (18%)	0 (0%)	27 (82%)	33 (100%)
matplotlib	3 (10%)	1 (3%)	0 (0%)	26 (87%)	30 (100%)
pandas	4 (2%)	0 (0%)	0 (0%)	165 (98%)	169 (100%)
sanic	0 (0%)	0 (0%)	0 (0%)	5 (100%)	5 (100%)
scrapy	0 (0%)	2 (5%)	0 (0%)	38 (95%)	40 (100%)
spacy	1 (10%)	0 (0%)	0 (0%)	9 (90%)	10 (100%)
thefuck	0 (0%)	0 (0%)	0 (0%)	32 (100%)	32 (100%)
tornado	0 (0%)	0 (0%)	0 (0%)	16 (100%)	16 (100%)
tqdm	0 (0%)	0 (0%)	0 (0%)	9 (100%)	9 (100%)
youtube-dl	0 (0%)	0 (0%)	0 (0%)	43 (100%)	43 (100%)
Total	12 (2%)	9 (2%)	3 (1%)	477 (95%)	501 (100%)

projects are sorted based on their running time in descending order, with the project "pandas" having the highest running time of 963 minutes, followed by "luigi," "scrapy," and so on.

Project	Running Time (minutes)
pandas	963
luigi	510
scrapy	268
keras	230
black	214
fastapi	197
thefuck	195
sanic	136
spacy	131
ansible	80
tqdm	36
youtube-dl	59
cookiecutter	40
httpie	39
matplotlib	26
tornado	14
pysnooper	5

IV. DISCUSSION

A. What makes reproduction easy?

The ease of bug reproduction in the BugsInPy dataset can be attributed to several factors:

1. **Automatation**: The `bugsinpy-testall` script provides an automated approach to reproducing and testing bugs in Python projects. The script streamlines the reproduction process, minimizes manual effort, and ensures we use a consistent procedure on each project. Note that these automation scripts must be carefully written and maintained. In particular, the original scripts did not have `set -e`, so some intermediate step might fail without alerting the user.
2. **Environment/Package manager**: The Conda environment/package manager simplifies the management of project dependencies. The crucial insight is that Conda can install packages in a local environment without interfering with global, system-wide packages. Conda makes it possible to define project-specific versions of libraries that a platform-specific system-wide package manager would normally manage.
3. **Lack of non-deterministic bugs**: None of the bugs in our dataset are race-conditions. Our scope is limited to constructing a reproducible software environment consistent with the original bug, and then the bug will show itself deterministically.

These factors collectively contribute to the ease of reproducing bugs in the BugsInPy dataset, providing a reliable and efficient framework for bug analysis and investigation.

B. What makes reproduction hard?

Despite the facilitative factors mentioned above, bug reproduction can still present challenges due to the following factors:

1. **Resource constraints during building**: The software environment can involve a computationally-expensive step of building software from source code. Reproducing and testing many bugs within limited resources may result in longer reproduction times and potential resource limitations. Our script creates many Conda environments. These can be expensive to store, and we cannot, for example, archive our environments in GitHub due to space constraints.
2. **Missing packages in Conda**: Unfortunately, not all Pip packages and versions exist in our selected Conda repositories.

Addressing these challenges requires careful consideration of project-specific factors and may involve additional research, debugging techniques, and resources to ensure accurate and reliable bug reproduction.

C. Recommendations to artifact authors

For authors providing Python research artifacts, the following recommendations can enhance the reproducibility of their artifacts:

1. **Make it automatic/easy to use**: The BugsInPy dataset has Python versions, but there is no automation to switch to a specific version, so users are unlikely to do so. Our improved version uses Conda to switch to the correct Python version automatically.
2. **`requirements.txt` is not enough**: Pip cannot handle library dependencies. Researchers should provide a container, a Conda lockfile, Spack lockfile, or other detailed environment specification.
3. **Archival storage**: Ensure that the artifact repository is archived in long-term storage, such as Zenodo or FigShare, so it does not bit rot.

D. Threats to Validity

Some of these bugs we find unreproducible *are* actually reproducible, but the error is on our side. Our efforts reflect an "average" user effort with limited resources, not a researcher with infinite time and resources.

Our work may not be reproducible for the following reasons:

1. Although we pin our Docker base images exact version, the source location (DockerHub) may stop hosting our image (e.g., goes out of business, ends free tier). In this case, one would need to change the base image, but it should still work, so long as that base image has Conda.
2. Conda package repositories can stop existing (e.g., if Anaconda goes out of business) or drop the old package versions we refer to. However, the definition of Conda packages describes how to build the packages from source code. The package definitions are smaller than the binaries, so they may remain longer.
3. The reproducers might need more computational resources to do the reproduction in a timely manner. However, we reuse Conda environments to minimize the computational cost. Furthermore, our scripts support reproducing just one project or just one bug from one project.

V. Conclusion

The study presented in this paper demonstrates the effectiveness of the BugsInPy dataset in reproducing and testing bugs in Python projects. The standardized and automated approach provided by the `bugsinpy-testall` script, coupled with the use of Conda for dependency management, streamlines the bug reproduction process and enhances its ease. The high success rate in reproducing bugs, with over 95% of bugs passing the tests, indicates the reliability and accuracy of the bug fixes in the dataset.

However, our experiments still depend on commercial organizations continuing to store software for free (GitHub, PyPI, Conda, Dockerhub). Challenges still exist in creating a truly long-term reproducible software environment.

The broader implications of the workflow used in the paper would be beneficial. The workflow of using standardized and automated approach, coupled with the use of Conda for dependency management, showcases the potential for improved bug reproduction processes. This workflow could not only be useful for BugsInPy but possibly also for other bug datasets to validate the reliability and accuracy of bug datasets in a more user friendly manner. However, challenges encountered, such as complex codebases and interdependencies, may vary across different projects and languages.

VI. Acknowledgments

References

[1] R. Widyasari, S. Q. Sim, C. Lok, *et al.*, "BugsInPy: A database of existing bugs in python programs to enable controlled testing and debugging studies," in *Proceedings of the 28th ACM Joint Meeting on European Software Engineering Conference and Symposium on the Foundations of Software Engineering*, ser. ESEC/FSE 2020, New York, NY, USA: Association for Computing Machinery, Nov. 8, 2020, pp. 1556–1560, ISBN: 978-1-4503-7043-1. DOI: 10.1145/3368089.3417943. [Online]. Available: https://doi.org/10.1145/3368089.3417943 (visited on 07/08/2023).

[2] S. Mukherjee, A. Almanza, and C. Rubio-González, "Fixing dependency errors for python build reproducibility," in *Proceedings of the 30th ACM SIGSOFT International Symposium on Software Testing and Analysis*, Virtual Denmark: ACM, Jul. 11, 2021, pp. 439–451, ISBN: 978-1-4503-8459-9. DOI: 10.1145/3460319.3464797. [Online]. Available: https://dl.acm.org/doi/10.1145/3460319.3464797 (visited on 07/17/2023).

[3] M. Smytzek and A. Zeller, "SFLKit: A workbench for statistical fault localization," in *Proceedings of the 30th ACM Joint European Software Engineering Conference and Symposium on the Foundations of Software Engineering*, Singapore Singapore: ACM, Nov. 7, 2022, pp. 1701–1705, ISBN: 978-1-4503-9413-0. DOI: 10.1145/3540250.3558915. [Online]. Available: https://dl.acm.org/doi/10.1145/3540250.3558915 (visited on 07/17/2023).

[4] T. Hirsch and B. Hofer, "A systematic literature review on benchmarks for evaluating debugging approaches," *Journal of Systems and Software*, vol. 192, p. 111 423, Oct. 2022, ISSN: 01641212. DOI: 10.1016/j.jss.2022.111423. [Online]. Available: https://linkinghub.elsevier.com/retrieve/pii/S0164121222001303 (visited on 07/17/2023).

[5] S. Lukasczyk, F. Kroiß, and G. Fraser, "An empirical study of automated unit test generation for python," *Empirical Software Engineering*, vol. 28, no. 2, p. 36, Mar. 2023, ISSN: 1382-3256, 1573-7616. DOI: 10.1007/s10664-022-10248-w. [Online]. Available: https://link.springer.com/10.1007/s10664-022-10248-w (visited on 07/17/2023).

[6] E. N. Akimova, A. Y. Bersenev, A. A. Deikov, *et al.*, "A survey on software defect prediction using deep learning," *Mathematics*, vol. 9, no. 11, p. 1180, May 24, 2021, ISSN: 2227-7390. DOI: 10.3390/math9111180. [Online]. Available: https://www.mdpi.com/2227-7390/9/11/1180 (visited on 07/17/2023).

[7] A. I. staff. "Artifact review and badging." (Aug. 24, 2020), [Online]. Available: https://www.acm.org/pub lications/policies/artifact-review-and-badging-current (visited on 01/19/2023).

[8] "Cmdoption-no-binary - pip install - pip documentation v23.2." (), [Online]. Available: https://pip.pypa.io/en/s table/cli/pip_install/#cmdoption-no-binary (visited on 07/17/2023).

[9] "Installation — matplotlib 3.7.2 documentation." (), [Online]. Available: https://matplotlib.org/stable/use rs/installing/index.html (visited on 07/17/2023).

[10] "Fixes -e option requires 1 argument. by faustinoaq · pull request #68 · soarsmu/BugsInPy," GitHub. (), [Online]. Available: https://github.com/soarsmu/BugsInPy/pull/68 (visited on 07/17/2023).

APPENDIX
CODE, DATA, AND REPRODUCING

A rolling release of all our code and data can be found at https://github.com/reproducing-research-projects/BugsInPy.

Our code includes:

- `Dockerfile` docker file setup to build projects images.
- `docker-compose.yml` docker orchestration to run projects containers.
- `framework/bin/bugsinpy-testall` script to automate execution of BugsInPy framework scripts.

To reproduce all bugs in a project, for example `httpie`, run:

```
$ rm -f projects/bugsinpy-index.csv
$ docker compose up setup httpie --build
Cleaning up temp folder ...
Reproducing bugs please wait ...
------------------------
httpie,1,buggy,fail
...
```

After these commands, the new results will be in the file named `bugsinpy-index.csv`.

See `README.md` for more detailed information.